

Adversarial Machine Learning

Arthea Valderrama
Faculty Advisor: Dr. Ian Chen
Committee Member: Dr. Sirong Lin

November 20, 2025

Contents

1 Abstract	2
2 Motivation	2
3 Attack Scenarios	3
3.1 Carlini & Wagner	3
3.2 Simple Black-box Attack	4
4 Anomaly Detection	4
5 Procedure	5
6 Results	5
7 Conclusion	7
A Appendix A: Model Creation	8
B Appendix B: Attack Implementation	12
C Appendix C: Anomaly Detection	25
D Appendix D: Web Application Deployment	29

1 Abstract

Machine Learning (ML) has become integral to daily life, due to its use in autonomous systems, the health-care industry, etc., because of its capability to produce a desired outcome based off of learning from training data. However, adversarial attacks attempt to deceive learning models by implementing a defective input, compromising the effectiveness as a result. We explore the mechanisms behind attacks on convolutional neural networks (CNNs). This project aims to contribute to the development of safer ML models to ensure their reliability, through evaluating the effectiveness of adversarial attacks.

2 Motivation

This is an addition to my current capstone project through UML's collaboration with Schneider Electric. The company is a multinational corporation specializing in energy management and automation. One of the company's ambitions is ensuring robust automation processes for increased efficiency and quality of services.

In an autonomous setting that requires quality assurance for real-time classification, it is important that a model is adaptable to various conditions to ensure quality decision making. In extension to the existing project's image classification system on household objects in an assembly line, I expanded the dataset for cancer detection in a healthcare setting to demonstrate the transferability of ML models and its attacks.

ML assisted detection of brain cancer through MRI imaging can save resources and improve patient outcomes through early detection and accurate patient prognosis estimation. However, an adversarial attack would lead to errors in tumor recognition, harming patient health and treatment. Such an attack could occur for a variety of reasons: undermining trust in critical systems, financial gain, etc. By identifying vulnerabilities in brain cancer imaging detection we can ensure the security and trustworthiness of Ai.

3 Attack Scenarios

3.1 Carlini & Wagner




Figure 1: Example implementation of C&W attack

The Carlini & Wagner (C&W)[1] attack finds the smallest perturbation to a model's input data, while also increasing the chances of model misclassification through optimizing an iterative algorithm. The algorithm is defined as: **minimize** $||\delta|| + c * f(x + \delta)$, **such that** $x + \delta \in [0, 1]^n$, where δ is a small change made to an image x that changes its classification, and c is a constant that advances the trade-off between the size of the change and effectiveness of the attack. It is a white-box attack, meaning the attacker can launch the attack if they have knowledge on the model's information, such as its architecture and its parameters.

3.2 Simple Black-box Attack




Figure 2: Small simBA distortion




Figure 3: Higher distortion of simBA

While the C&W attack is a white-box attack that requires the adversary to have full knowledge of the model it is attacking, the Simple Black-box Attack (simBA)[\[2\]](#) is a type of black-box attack that does not need complete information about the model. In that case, simBA is more applicable in many scenarios. SimBA works by performing a random search on the pixels of an image to find minuscule perturbations in order to change the image away from its predicted classification and continuously checking the probability of success in changing input.

4 Anomaly Detection

An anomaly detection [\[3\]](#) model is a specialized CNN that identifies data points that significantly deviate from "expected" within the dataset. The model works by taking in a binary input: the normal patterns within the dataset and anomalous samples that deviate from the learned patterns. By incorporating aberrant samples in its training, the model is then able to flag unusual input. This can protect an AI-based system in real-time from instances that could result in harmful losses. Other use cases of anomaly detection within AI applications are detecting fraudulent financial transactions and suspicious network traffic in cybersecurity.

5 Procedure

A crucial factor in ML training is the selection of a robust dataset. I obtained a brain tumor MRI dataset of 7,022 images[4], and split the dataset 80% to train the model and 20% to evaluate the model's accuracy. I chose to select a large dataset for improved generalization with more diverse data samples and to reduce overfitting. The dataset includes 4 classifications of tumors: no tumor, glioma (tumors in the glial cells in the brain or spinal cord), meningioma (tumors that develops in the meninges), and pituitary (tumors in the pituitary gland).

In accordance with my capstone project, a convolutional neural network (CNN) was the model architecture chosen, specifically in the form of a DenseNet model. This model architecture consists of dense layers where each layer is connected to its previous layers in a feed-forward fashion. The DenseNet model is helpful for image classification, as its compact network improves information propagation through the network, has an enhanced ability to learn complex features, and requires less parameters in deep whilst not needing much computation.

The main programming language used in this project is Python. Python is ideal for working with ML, as it has numerous libraries for data pre-processing and evaluation, such as Sci-kit Learn, Numpy, Pandas, and Matplotlib. The main library used for ML is PyTorch, which allowed me to import a wide range of pretrained Neural Networks to improve model accuracy and functionality.

When coding my training script in Python to train the model, I created a class to preprocess the MRI images into classes for the model to distinguish between, and another class to define the model's parameters and functions. The model's class consists of a initialization function that uses features from Pytorch's built-in pretrained DenseNet model, a forward function to define the model's learning and optimization, and a training function to define the number of iterations the model passes through its learning algorithm while tracking its loss. Utilizing a pretrained model improved the model's learning rate and testing accuracy on samples it has not seen before.

Creating the attack code in Python, I modified widely available C&W and simBA attack code to apply to my MRI imaging model. Thus, helping to demonstrate the generalizability of the attack across different imaging models. I applied the attacks on the original dataset of 7,022 MRI images to generate enough adversarial samples for the anomaly detection model.

With the anomaly detection model, I then combined both the original training data and adversarial data. Thus, in total it learned the patterns and behavior between 14,044 MRI images to recognize between suspicious outliers and regular input.

Lastly, to create a simple web app to test the models on a user-friendly interface. Through the web app, I used the Flask library, HTML, CSS, and JavaScript code for live predictions on the trained dataset. I then incorporated Google's AI Gemini to detail more information about adversarial attacks, information about diagnostics, and educate the implications of the model's findings. A chat-bot was added to allow the user to talk more with the AI for any follow-up questions.

6 Results

Through testing over a sample batch size of 32 images, I was able to calculate the attack success rate of the C&W, simBA, and Anomaly Detection model. The C&W attack is more successful than simBA by a




Figure 4: Success Rate of Attacks and Anomaly Detections

factor of around 20%. The success rate of the anomaly detection method on the C&W attack proves the effectiveness further, as the model has detected less C&W samples than the simBA method. With a higher




Figure 5: Success Rate of Attacks and Anomaly Detections

δ , the success rates of attacks tend to be higher but cause more visual distortions. The C&W generally does not need a high δ to be effective, although simBA requires a factor of 1.0+ to start becoming successful in fooling the model. Another comparison that can be made between the two models is that the C&W, while more effective, requires much more time than the simBA attack to deploy on average.




Figure 6: Time taken to deploy the attack

7 Conclusion

As AI systems become increasingly integrated into society, it is very important that security considerations are taken when deploying and creating AI systems. Models often influence critical decisions, and vulnerabilities that exploit these models can have potentially devastating consequences in quality of life, finances, etc. The example adversarial attacks of C&W and simBA showcase the generalizability and adaptability of these attacks across different use cases. These attacks can have devastating consequences on industries and organizations that rely on computer vision or image classification. Without strong security measures and further research into the possibilities of attacks on different types of AI models, the risk of model exploitation could harm societal trust in AI technology. Thus, ensuring the security of AI is essential to safeguard against manipulation, and uphold the reliability of automated decision-making.

A Appendix A: Model Creation

The following code document details the creation of an image classification model, originally imported from my senior capstone project and adapted to be utilized for MRI imaging purposes.

```
1 import torch
2 import torchvision.transforms as transforms
3 from torch.utils.data import Dataset, DataLoader
4 import os
5 from torch import nn
6 import torch.nn.functional as F
7 from torchvision import models
8 import matplotlib.pyplot as plt
9 from PIL import Image
10 import numpy as np
11 from tqdm import tqdm
12
13 class CustomDataset(Dataset):
14     def __init__(self, root_dir, transform=None):
15         self.root_dir = root_dir
16         self.transform = transform
17         self.classes = ['glioma', 'meningioma', 'notumor', 'pituitary']
18         self.image_paths = []
19         self.labels = []
20
21         for idx, class_name in enumerate(self.classes):
22             class_dir = os.path.join(root_dir, class_name)
23             for image_name in os.listdir(class_dir):
24                 if image_name.endswith(".jpg") or image_name.endswith
25 ("png"):
26                     self.image_paths.append(os.path.join(class_dir,
27 image_name))
28                     self.labels.append(idx)
29
30     def __len__(self):
31         return len(self.image_paths)
32
33     def __getitem__(self, idx):
34         img_name = self.image_paths[idx]
35         image = Image.open(img_name).convert('RGB')
36         label = self.labels[idx]
37
38         if self.transform:
39             image = self.transform(image)
40
41         resize_transform = transforms.Compose([
42             transforms.Resize(256),
43             transforms.CenterCrop(224),
44             transforms.ToTensor(),
45             transforms.Normalize([0.485, 0.45, 0.406], [0.229, 0.224, 0.225])]
```

```

42     transforms.Resize((64, 64)),
43     transforms.ToTensor()
44 ])
45
46 class_names = ['glioma', 'meningioma', 'notumor', 'pituitary']
47 class_id_to_name = {i: class_name for i, class_name in enumerate(
48     class_names)}
49 class_to_idx = {class_name: i for i, class_name in enumerate(
50     class_names)}
51
52 class FineTunedDenseNet(nn.Module):
53     def __init__(self, pretrained_model, num_classes=4):
54         super(FineTunedDenseNet, self).__init__()
55         self.features = pretrained_model.features
56         self.classifier = nn.Linear(pretrained_model.classifier.
57             in_features, num_classes)
58
59     def forward(self, x):
60         x = self.features(x)
61         x = F.adaptive_avg_pool2d(x, (1, 1)).view(x.size(0), -1)
62         x = self.classifier(x)
63         return x
64
65     def train_model(self, model_name, input_size, epochs,
66     learning_rate, train_loader, test_loader, device):
67         criterion = nn.CrossEntropyLoss()
68         optimizer = torch.optim.Adam(self.parameters(), lr=
69         learning_rate)
70         scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
71         step_size=5, gamma=0.1)
72
73         best_accuracy = 0
74         train_loss_values = []
75         test_accuracy_values = []
76
77         for epoch in range(epochs):
78             self.train()
79             running_loss = 0.0
80             correct = 0
81             total = 0
82             progress_bar = tqdm(train_loader, desc=f"Epoch {epoch+1}/{epochs}",
83             leave=False, ncols=100)
84
85             for i, (images, labels) in enumerate(progress_bar):
86                 images, labels = images.to(device), labels.to(device)
87
88                 optimizer.zero_grad()
89                 outputs = self(images)

```

```

83         loss = criterion(outputs, labels)
84         loss.backward()
85         optimizer.step()
86
87         running_loss += loss.item()
88         _, predicted = torch.max(outputs, 1)
89         total += labels.size(0)
90         correct += (predicted == labels).sum().item()
91
92         progress_bar.set_postfix(loss=running_loss / (i + 1),
93 accuracy=100 * correct / total)
94
95         test_accuracy = 100 * correct / total
96         if test_accuracy > best_accuracy:
97             best_accuracy = test_accuracy
98             torch.save(self.state_dict(), "best_densenet_model.pth"
99         ")
100
101         scheduler.step()
102         train_loss_values.append(running_loss / len(train_loader))
103         test_accuracy_values.append(test_accuracy)
104
105         self.plot_loss(model_name, input_size, learning_rate,
106 train_loss_values, test_accuracy_values)
107
108     def plot_loss(self, model_name, input_size, learning_rate,
109 train_loss_values, test_accuracy_values):
110         plt.figure(figsize=(10, 5))
111         plt.plot(train_loss_values, label="Training Loss", color='blue'
112     )
113         plt.xlabel("Training Steps")
114         plt.ylabel("Loss")
115         plt.title("Training Loss Over Time")
116         plt.grid(True)
117         plt.legend()
118         plt.savefig(f"training_loss_{model_name}_{input_size}_lr{learning_rate}.png")
119         plt.show()
120
121         plt.figure(figsize=(10, 5))
122         plt.plot(test_accuracy_values, label="Test Accuracy", color='orange')
123         plt.xlabel("Steps")
124         plt.ylabel("Accuracy (%)")
125         plt.title("Test Accuracy Over Time")
126         plt.grid(True)
127         plt.legend()

```

```

123     plt.savefig(f"test_accuracy_{model_name}_{input_size}_lr{learning_rate}.png")
124     plt.show()
125
126     def save(self, model_file):
127         torch.save(self.state_dict(), model_file)
128
129 if __name__ == "__main__":
130     root_dir = '/Users/arthe/honors/brain/training'
131     model_name = 'FineTunedDenseNet'
132     batch_size = 32
133     learning_rate = 0.0003
134     epochs = 10
135     input_size = 64
136
137     image_transforms = transforms.Compose([
138         transforms.RandomResizedCrop(input_size),
139         transforms.RandomHorizontalFlip(),
140         transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.2),
141         transforms.ToTensor(),
142         transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
143     ])
144
145     full_dataset = CustomDataset(root_dir, transform=image_transforms)
146     train_size = int(0.8 * len(full_dataset))
147     test_size = len(full_dataset) - train_size
148     train_dataset, test_dataset = torch.utils.data.random_split(
149         full_dataset, [train_size, test_size])
150     train_loader = DataLoader(train_dataset, batch_size=batch_size,
151     shuffle=True, num_workers=4)
152     test_loader = DataLoader(test_dataset, batch_size=batch_size,
153     shuffle=False, num_workers=4)
154
155     print(f"Training {model_name}")
156     device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
157     print(f"Using device {device}")
158
159     pretrained_densenet = models.densenet121(pretrained=True)
160     model = FineTunedDenseNet(pretrained_model=pretrained_densenet,
161     num_classes=4)
162     model.to(device)
163
164     model.train_model(model_name, input_size, epochs, learning_rate,
165     train_loader, test_loader, device=device)

```

```

162 model.to('cpu')
163 model_path = f'{model_name}_{input_size}_lr{learning_rate}.pt'
164 torch.save(model.state_dict(), model_path)
165 print(f"Model saved at {model_path}")

```

Listing 1: train.py

B Appendix B: Attack Implementation

This section details the Python implementation of the C&W[5] attack and the simBA[6] attack, originally sourced from GitHub and then modified to suit the image classification model being used in this project. The first document is my implementation of the C&W attack, with additional functions to save adversarial samples and display the original classification and modded image for viewing purposes. The next lines of code is my simBA implementation, with similar functions to save adversarial samples and display the difference between the original image and the modded image.

```

1 import os
2 import torch
3 import time
4 import torch.nn.functional as F
5 import torch.optim as optim
6 from torchvision import models
7 from torch import nn
8 from torchvision import transforms
9 import matplotlib.pyplot as plt
10 from PIL import Image
11
12 from train import class_names, class_id_to_name, class_to_idx,
13     CustomDataset
14 from train import FineTunedDenseNet
15
16 model_file = "/Users/arthe/honors/FineTunedDenseNet_64_lr0.0003.pt"
17
18 script_dir = os.path.dirname(os.path.abspath(__file__))
19 adversarial_save_dir = os.path.join(script_dir, "adversarial_images1")
20 os.makedirs(adversarial_save_dir, exist_ok=True)
21
22 resize_transform = transforms.Resize((224, 224), interpolation=
23     transforms.InterpolationMode.BILINEAR)
24
25 full_transform = transforms.Compose([
26     resize_transform,
27     transforms.ToTensor()
28 ])
29
30 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
31 pretrained_densenet = models.densenet121(pretrained=True)

```



```

71         one_hot_labels = torch.eye(outputs.shape[1], device=device)
72     [labels]
73         i, _ = torch.max((1 - one_hot_labels) * outputs, dim=1)
74         j = torch.masked_select(outputs, one_hot_labels.bool())
75         return torch.clamp(j - i, min=-kappa)
76
77
78     w = torch.zeros_like(images, requires_grad=True).to(device)
79     optimizer = optim.Adam([w], lr=learning_rate)
80     prev = 1e10
81
82
83     for step in range(max_iter):
84         a = 0.5 * (torch.tanh(w) + 1)
85
86         loss1 = F.mse_loss(a, images, reduction='sum')
87         loss2 = torch.sum(c * f(a))
88         cost = loss1 + loss2
89
90         optimizer.zero_grad()
91         cost.backward()
92         optimizer.step()
93
94         if step % (max_iter // 10) == 0:
95             if cost > prev:
96                 print("Attack Stopped due to CONVERGENCE....")
97                 return a
98             prev = cost
99
100            if (step + 1) % 100 == 0 or step == 0:
101                print(f"- Learning Progress: {(step + 1) / max_iter * 100:.2f}%", end="\r")
102
103    attack_images = 0.5 * (torch.tanh(w) + 1)
104    return attack_images
105
106
107 def save_image(image, true_label, predicted_label, original_size, index,
108 , prefix="adversarial_example"):
109
110     image = transforms.ToPILImage()(image.cpu().detach())
111
112     image_resized = image.resize(original_size, Image.LANCZOS)
113
114     os.makedirs(adversarial_save_dir, exist_ok=True)
115
116     file_name = os.path.join(
117         adversarial_save_dir,
118         f"{prefix}_{index}_true_{true_label}_pred_{predicted_label}.png"
119     )
120
121
122
123
124

```

```

115     image_resized.save(file_name)
116
117
118 def display_images(original_image, adversarial_image, true_label,
119                     predicted_label):
120
121     fig, axes = plt.subplots(1, 2, figsize=(12, 6))
122
123     axes[0].imshow(original_image.permute(1, 2, 0).cpu().detach().numpy())
124     axes[0].set_title(f"Original: {true_label}")
125     axes[0].axis('off')
126
127     axes[1].imshow(adversarial_image.permute(1, 2, 0).cpu().detach().numpy())
128     axes[1].set_title(f"Adversarial: {predicted_label}")
129     axes[1].axis('off')
130
131     plt.show()
132
133 if __name__ == "__main__":
134     import argparse
135
136     targeted = False
137     target_class = "glioma"
138     c = 20
139     kappa = 0
140     learning_rate = 0.01
141     max_iter = 1000
142
143     data_folder = "/Users/arthe/honors/brain/training"
144     full_dataset = CustomDataset(data_folder, transform=full_transform)
145     train_size = int(0.8 * len(full_dataset))
146     test_size = len(full_dataset) - train_size
147     train_dataset, test_dataset = torch.utils.data.random_split(
148         full_dataset, [train_size, test_size])
149
150     batch_size = 1
151     test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size, shuffle=False)
152
153     attack_accuracy = 0
154     correct = 0
155     total = 0
156
157     for images, labels in test_loader:
158         images, labels = images.to(device), labels.to(device)

```

```

158     original_size = images.shape[2], images.shape[3]
159
160     start_time = time.time()
161
162     if targeted:
163         if target_class is None:
164             raise ValueError("Target class must be specified for
targeted attacks.")
165         try:
166             target_class_idx = class_to_idx[target_class]
167         except KeyError:
168             raise ValueError(f"Target class '{target_class}' not
found in class_to_idx.")
169         target_labels = torch.full_like(labels, target_class_idx)
170         adversarial_images = cw_l2_attack(
171             model, images, labels, target_labels=target_labels,
172             targeted=True, c=c,
173             kappa=kappa, learning_rate=learning_rate, max_iter=
max_iter
174         )
175     else:
176         adversarial_images = cw_l2_attack(
177             model, images, labels, targeted=False, c=c,
178             kappa=kappa, learning_rate=learning_rate, max_iter=
max_iter
179         )
180
181         end_time = time.time()
182         attack_time = end_time - start_time
183         print(f"Time taken for attack: {attack_time:.2f} seconds")
184
185         adversarial_images_resized = resize_transform(
186             adversarial_images)
187
188         outputs = model(adversarial_images_resized)
189         _, predicted = torch.max(outputs, 1)
190
191         if targeted:
192             batch_correct = (predicted == target_labels).sum().item()
193         else:
194             batch_correct = (predicted != labels).sum().item()
195
196         batch_total = labels.size(0)
197         attack_accuracy += 100 * batch_correct / batch_total
198         correct += batch_correct
199         total += batch_total

```

```

199     print(f"Attack Accuracy on the batch: {100 * batch_correct / batch_total:.2f}%")
200
201     for i in range(len(adversarial_images)):
202         true_label = class_names[labels[i].item()]
203         predicted_label = class_names[predicted[i].item()]
204
205         save_image(adversarial_images[i], true_label,
206         predicted_label, original_size, i)
207
208         display_images(images[i], adversarial_images[i], true_label,
209         predicted_label)
210
211     overall_attack_accuracy = attack_accuracy / len(test_loader)
212     print(f"\nOverall Attack Accuracy: {overall_attack_accuracy:.2f}%")

```

Listing 2: attack.py

```

1 import torch
2 import utils
3 import torch.nn.functional as F
4 import torchvision.models as models
5 from torchvision.models import densenet121
6 import torchvision.transforms as transforms
7 from torchvision.utils import make_grid
8 import matplotlib.pyplot as plt
9 import time
10 from torchvision.datasets import ImageFolder
11 from torch.utils.data import DataLoader
12
13
14 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
15
16 class SimBA:
17
18     def __init__(self, model, dataset, image_size):
19         self.model = model.to(device)
20         self.dataset = dataset
21         self.image_size = image_size
22         self.model.eval()
23
24     def expand_vector(self, x, size):
25         batch_size = x.size(0)
26         x = x.view(-1, 3, size, size)
27         z = torch.zeros(batch_size, 3, self.image_size, self.image_size
28     )
28         z[:, :, :size, :size] = x
29         return z
30

```

```

31     def normalize(self, x):
32         return utils.apply_normalization(x.to(device), self.dataset)
33
34
35     def get_probs(self, x, y):
36         output = self.model(self.normalize(x)).cpu()
37         probs = torch.index_select(F.softmax(output, dim=-1).data, 1, y
38 )
39         return torch.diag(probs)
40
41     def get_preds(self, x):
42         output = self.model(self.normalize(x)).cpu()
43         _, preds = output.data.max(1)
44         return preds
45
46     def simba_single(self, x, y, num_iters=10000, epsilon=2.5, targeted
47 =False):
48         n_dims = x.view(1, -1).size(1)
49         perm = torch.randperm(n_dims)
50         last_prob = self.get_probs(x, y)
51         for i in range(num_iters):
52             diff = torch.zeros(n_dims)
53             diff[perm[i]] = epsilon
54             left_prob = self.get_probs((x - diff.view(x.size())).clamp
55 (0, 1), y)
56             if targeted != (left_prob < last_prob):
57                 x = (x - diff.view(x.size())).clamp(0, 1)
58                 last_prob = left_prob
59             else:
60                 right_prob = self.get_probs((x + diff.view(x.size())).c
61 clamp(0, 1), y)
62                 if targeted != (right_prob < last_prob):
63                     x = (x + diff.view(x.size())).clamp(0, 1)
64                     last_prob = right_prob
65             if i % 10 == 0:
66                 print(last_prob)
67         return x.squeeze()
68
69
70     def simba_batch(self, images_batch, labels_batch, max_iters,
71 freq_dims, stride, epsilon, linf_bound=0.0,
72 order='rand', targeted=False, pixel_attack=False,
73 log_every=1):
74         batch_size = images_batch.size(0)
75         image_size = images_batch.size(2)
76         assert self.image_size == image_size
77         if order == 'rand':
78             indices = torch.randperm(3 * freq_dims * freq_dims) [:
```

```

max_iters]
    elif order == 'diag':
        indices = utils.diagonal_order(image_size, 3) [:max_iters]
    elif order == 'strided':
        indices = utils.block_order(image_size, 3, initial_size=
freq_dims, stride=stride) [:max_iters]
    else:
        indices = utils.block_order(image_size, 3) [:max_iters]
if order == 'rand':
    expand_dims = freq_dims
else:
    expand_dims = image_size
n_dims = 3 * expand_dims * expand_dims
x = torch.zeros(batch_size, n_dims)
probs = torch.zeros(batch_size, max_iters)
succs = torch.zeros(batch_size, max_iters)
queries = torch.zeros(batch_size, max_iters)
l2_norms = torch.zeros(batch_size, max_iters)
linf_norms = torch.zeros(batch_size, max_iters)
prev_probs = self.get_probs(images_batch, labels_batch)
preds = self.get_preds(images_batch)
if pixel_attack:
    trans = lambda z: z
else:
    trans = lambda z: utils.block_idct(z, block_size=image_size
, linf_bound=linf_bound)
remaining_indices = torch.arange(0, batch_size).long()
for k in range(max_iters):
    dim = indices[k]
    expanded = (images_batch[remaining_indices] + trans(self.
expand_vector(x[remaining_indices], expand_dims))).clamp(0, 1)
    perturbation = trans(self.expand_vector(x, expand_dims))
    l2_norms[:, k] = perturbation.view(batch_size, -1).norm(2,
1)
    linf_norms[:, k] = perturbation.view(batch_size, -1).abs().max(1)[0]
    preds_next = self.get_preds(expanded)
    preds[remaining_indices] = preds_next
    if targeted:
        remaining = preds.ne(labels_batch)
    else:
        remaining = preds.eq(labels_batch)
    if remaining.sum() == 0:
        adv = (images_batch + trans(self.expand_vector(x,
expand_dims))).clamp(0, 1)
        probs_k = self.get_probs(adv, labels_batch)
        probs[:, k:] = probs_k.unsqueeze(1).repeat(1, max_iters
- k)

```

```

113         succs[:, k:] = torch.ones(batch_size, max_iters - k)
114         queries[:, k:] = torch.zeros(batch_size, max_iters - k)
115         break
116     remaining_indices = torch.arange(0, batch_size)[remaining].
117     long()
118     if k > 0:
119         succs[:, k] = ~remaining
120         diff = torch.zeros(remaining.sum(), n_dims)
121         diff[:, dim] = epsilon
122         left_vec = x[remaining_indices] - diff
123         right_vec = x[remaining_indices] + diff
124         adv = (images_batch[remaining_indices] + trans(self.
125         expand_vector(left_vec, expand_dims))).clamp(0, 1)
126         left_probs = self.get_probs(adv, labels_batch[
127         remaining_indices])
128         queries_k = torch.zeros(batch_size)
129         queries_k[remaining_indices] += 1
130         if targeted:
131             improved = left_probs.gt(prev_probs[remaining_indices])
132         else:
133             improved = left_probs.lt(prev_probs[remaining_indices])
134         if improved.sum() < remaining_indices.size(0):
135             queries_k[remaining_indices[~improved]] += 1
136             adv = (images_batch[remaining_indices] + trans(self.
137             expand_vector(right_vec, expand_dims))).clamp(0, 1)
138             right_probs = self.get_probs(adv, labels_batch[
139             remaining_indices])
140             if targeted:
141                 right_improved = right_probs.gt(torch.max(prev_probs[
142                 remaining_indices], left_probs))
143             else:
144                 right_improved = right_probs.lt(torch.min(prev_probs[
145                 remaining_indices], left_probs))
146             probs_k = prev_probs.clone()
147             if improved.sum() > 0:
148                 left_indices = remaining_indices[improved]
149                 left_mask_remaining = improved.unsqueeze(1).repeat(1,
n_dims)
x[left_indices] = left_vec[left_mask_remaining].view
(-1, n_dims)
probs_k[left_indices] = left_probs[improved]
if right_improved.sum() > 0:
    right_indices = remaining_indices[right_improved]
    right_mask_remaining = right_improved.unsqueeze(1).
repeat(1, n_dims)
x[right_indices] = right_vec[right_mask_remaining].view
(-1, n_dims)
probs_k[right_indices] = right_probs[right_improved]

```

```

150     probs[:, k] = probs_k
151     queries[:, k] = queries_k
152     prev_probs = probs[:, k]
153     if (k + 1) % log_every == 0 or k == max_iters - 1:
154         print('Iteration %d: queries = %.4f, prob = %.4f,' %
155             k + 1, queries.sum(1).mean(), probs[:, k].mean()
156             (), remaining.float().mean()))
157     expanded = (images_batch + trans(self.expand_vector(x,
158                                         expand_dims))).clamp(0, 1)
159     preds = self.get_preds(expanded)
160     if targeted:
161         remaining = preds.ne(labels_batch)
162     else:
163         remaining = preds.eq(labels_batch)
164     succs[:, max_iters-1] = ~remaining
165     return expanded, probs, succs, queries, l2_norms, linf_norms
166 transform = transforms.Compose([
167     transforms.Resize((224, 224)),
168     transforms.ToTensor(),
169 ])
170
171 dataset_path = '/Users/arthe/honors/brain/training'
172 dataset = ImageFolder(root=dataset_path, transform=transform)
173 dataloader = DataLoader(dataset, batch_size=2, shuffle=True)
174
175 model = densenet121(pretrained=True)
176 model.classifier = torch.nn.Linear(in_features=1024, out_features=4)
177
178 model.load_state_dict(torch.load('/Users/arthe/honors/
179     FineTunedDenseNet_64_lr0.0003.pt', map_location=device))
180 model.eval()
181
182 simba = SimBA(model=model, dataset='imagenet', image_size=224) #was 224
183
184 def simba_single_attack(simba, dataloader):
185     images_batch, labels_batch = next(iter(dataloader))
186     image = images_batch[0:1].to(device)
187     label = labels_batch[0:1].to(device)
188
189     class_name = dataloader.dataset.classes[label.item()]
190
191     start_time = time.time()
192     adv_image = simba.simba_single(image, label, num_iters=100, epsilon
193         =0.2, targeted=False)
194     elapsed_time = time.time() - start_time

```

```

193
194     adv_pred = simba.get_preds(adv_image.unsqueeze(0))
195     adv_class_name = dataloader.dataset.classes[adv_pred.item()]
196
197     plt.figure(figsize=(10, 5))
198
199     plt.subplot(1, 2, 1)
200     plt.imshow(image[0].permute(1, 2, 0).cpu().numpy())
201     plt.title(f"Original: {class_name} ({label.item()})")
202     plt.axis('off')
203
204     plt.subplot(1, 2, 2)
205     plt.imshow(adv_image.permute(1, 2, 0).cpu().numpy())
206     plt.title(f"Adversarial: {adv_class_name} ({adv_pred.item()})")
207     plt.axis('off')
208
209     plt.suptitle(f"SimBA Single Attack | Time: {elapsed_time:.2f}s")
210     plt.show()
211
212
213
214 def simba_full_dataset_attack(simba, dataloader, max_iters=200, epsilon
215     =0.2, freq_dims=21, stride=1, targeted=False):
216     total_images = 0
217     total_fooled = 0
218     total_time = 0.0
219
220     for images_batch, labels_batch in dataloader:
221         images_batch = images_batch.to(device)
222         labels_batch = labels_batch.to(device)
223
224         start_time = time.time()
225         adv_images, probs, succs, queries, l2_norms, linf_norms = simba
226             .simba_batch(
227                 images_batch=images_batch,
228                 labels_batch=labels_batch,
229                 max_iters=max_iters,
230                 freq_dims=freq_dims,
231                 stride=stride,
232                 epsilon=epsilon,
233                 linf_bound=0.0,
234                 order='rand',
235                 targeted=targeted,
236                 pixel_attack=False,
237                 log_every=50
238             )
239         elapsed_time = time.time() - start_time
240         total_time += elapsed_time

```

```

239     adv_preds = simba.get_preds(adv_images)
240     fooled = (adv_preds != labels_batch).sum().item()
241     total_fooled += fooled
242     total_images += len(labels_batch)
243
244     print(f"Batch: {total_images}/{len(dataloader.dataset)} -"
245           f"Fooled: {fooled}/{len(labels_batch)} - Time: {elapsed_time:.2f}s")
246
247     success_rate = (total_fooled / total_images) * 100
248     avg_time = total_time / len(dataloader)
249
250     print(f"\nSimBA Attack Completed")
251     print(f"Total Images: {total_images}")
252     print(f"Total Fooled: {total_fooled}")
253     print(f"Attack Success Rate: {success_rate:.2f}%")
254     print(f"Average Time per Batch: {avg_time:.2f}s")
255
256     return success_rate
257
258
259
260 def evaluate_adversarial_attacks(model, dataloader, targeted=False,
261                                   max_iters=200, freq_dims=21, stride=1, epsilon=2.5):
262
263     model.eval()
264     attack_accuracy = 0
265     correct = 0
266     total = 0
267
268     for images, labels in dataloader:
269         images, labels = images.to(device), labels.to(device)
270
271         start_time = time.time()
272
273         if targeted:
274             if target_class is None:
275                 raise ValueError("Target class must be specified for"
276                                  "targeted attacks.")
277             try:
278                 target_class_idx = dataloader.dataset.class_to_idx[
279                     target_class]
280             except KeyError:
281                 raise ValueError(f"Target class '{target_class}' not"
282                                 "found in class_to_idx.")
283             target_labels = torch.full_like(labels, target_class_idx)
284             adv_images, *_ = simba.simba_batch(
285                 images_batch=images,

```

```

282         labels_batch=labels,
283         max_iters=max_iters,
284         freq_dims=freq_dims,
285         stride=stride,
286         epsilon=epsilon,
287         targeted=True
288     )
289 else:
290     adv_images, *_ = simba.simba_batch(
291         images_batch=images,
292         labels_batch=labels,
293         max_iters=max_iters,
294         freq_dims=freq_dims,
295         stride=stride,
296         epsilon=epsilon,
297         targeted=False
298     )
299
300     end_time = time.time()
301     attack_time = end_time - start_time
302     print(f"Time taken for attack: {attack_time:.2f} seconds")
303
304     outputs = model(adv_images)
305     _, predicted = torch.max(outputs, 1)
306
307     if targeted:
308         batch_correct = (predicted == target_labels).sum().item()
309     else:
310         batch_correct = (predicted != labels).sum().item()
311
312     batch_total = labels.size(0)
313     attack_accuracy += 100 * batch_correct / batch_total
314     correct += batch_correct
315     total += batch_total
316
317     print(f"Attack Accuracy on the batch: {100 * batch_correct / batch_total:.2f}%")
318
319     overall_attack_accuracy = attack_accuracy / len(dataloader)
320     print(f"\nOverall Attack Accuracy: {overall_attack_accuracy:.2f}%")
321
322     return overall_attack_accuracy
323
324 #simba_single_attack(simba, dataloader)
325
326 #success_rate = simba_full_dataset_attack(simba, dataloader, max_iters
327 #                                         =200, epsilon=0.2)
327 overall_accuracy = evaluate_adversarial_attacks(model, dataloader,

```

```
targeted=False, max_iters=200)
```

Listing 3: simba.py

C Appendix C: Anomaly Detection

The next document of code is the creation of the anomaly detection model, which utilizes the attack implementation and unmodified data samples to train.

```
1 import torch
2 from train import FineTunedDenseNet
3 from torchvision import transforms, datasets
4 from torch.utils.data import DataLoader, ConcatDataset, Subset
5 from PIL import Image
6 from tqdm import tqdm
7 import matplotlib.pyplot as plt
8 from torchvision import models
9 import os
10 from sklearn.model_selection import train_test_split
11
12 original_data_dir = '/Users/arthe/honors/train'
13 adversarial_data_dir = '/Users/arthe/honors/adversarial_images'
14
15 transform = transforms.Compose([
16     transforms.Resize((224, 224)),
17     transforms.ToTensor(),
18 ])
19
20 original_dataset = datasets.ImageFolder(
21     root=original_data_dir,
22     transform=transform
23 )
24
25 adversarial_dataset = datasets.ImageFolder(
26     root=adversarial_data_dir,
27     transform=transform
28 )
29
30 class RelabeledDataset(torch.utils.data.Dataset):
31     def __init__(self, dataset, new_label):
32         self.dataset = dataset
33         self.new_label = new_label
34
35     def __len__(self):
36         return len(self.dataset)
37
38     def __getitem__(self, idx):
```

```

39         img, _ = self.dataset[idx]
40         return img, self.new_label
41
42 original_binary = RelabeledDataset(original_dataset, 0)
43 adversarial_binary = RelabeledDataset(adversarial_dataset, 1)
44 combined_dataset = ConcatDataset([original_binary, adversarial_binary])
45
46 all_data = [(img, label) for img, label in combined_dataset]
47 labels = [label for _, label in all_data]
48 train_idx, val_idx = train_test_split(
49     list(range(len(all_data))), test_size=0.2, stratify=labels,
50     random_state=42
51 )
52 train_dataset = Subset(combined_dataset, train_idx)
53 val_dataset = Subset(combined_dataset, val_idx)
54
55 class AnomalyDenseNet(torch.nn.Module):
56     def __init__(self, base_model: FineTunedDenseNet):
57         super(AnomalyDenseNet, self).__init__()
58         self.features = base_model.features
59
60         self.classifier = torch.nn.Sequential(
61             torch.nn.AdaptiveAvgPool2d((1, 1)),
62             torch.nn.Flatten(),
63             torch.nn.Linear(1024, 1)
64         )
65
66     def forward(self, x):
67         x = self.features(x)
68         out = self.classifier(x)
69         return out.squeeze()
70
71     def train_model(self, model_name, input_size, epochs, learning_rate,
72 , train_loader, val_loader, device):
73         optimizer = torch.optim.Adam(self.parameters(), lr=
learning_rate)
74         criterion = torch.nn.BCEWithLogitsLoss()
75
76         for epoch in range(epochs):
77             self.train()
78             running_loss, correct_train, total_train = 0.0, 0, 0
79             train_loader_tqdm = tqdm(train_loader, desc=f"Epoch {epoch
+1}/{epochs} [Train]", leave=False)
80
81                 for inputs, labels in train_loader_tqdm:
82                     inputs, labels = inputs.to(device), labels.float().to(
device)

```

```

82         optimizer.zero_grad()
83         outputs = self(inputs).view(-1)
84         labels = labels.view(-1)
85         loss = criterion(outputs, labels)
86         loss.backward()
87         optimizer.step()

88
89         batch_size = inputs.size(0)
90         running_loss += loss.item() * batch_size
91         preds = (torch.sigmoid(outputs) > 0.5).float()
92         correct_train += (preds == labels).sum().item()
93         total_train += batch_size

94
95         train_loader_tqdm.set_postfix({
96             'loss': f'{running_loss/total_train:.4f}',
97             'acc': f'{correct_train/total_train:.4f}'
98         })

99
100        epoch_loss = running_loss / len(train_loader.dataset)
101        epoch_acc = correct_train / total_train

102
103        self.eval()
104        val_loss, correct_val, total_val = 0.0, 0, 0
105        val_loader_tqdm = tqdm(val_loader, desc=f"Epoch {epoch+1}/{epochs} [Val]", leave=False)

106
107        with torch.no_grad():
108            for inputs, labels in val_loader_tqdm:
109                inputs, labels = inputs.to(device), labels.float().
110                to(device)
111                outputs = self(inputs).view(-1)
112                loss = criterion(outputs, labels)
113                batch_size = inputs.size(0)
114                val_loss += loss.item() * batch_size
115                preds = (torch.sigmoid(outputs) > 0.5).float()
116                correct_val += (preds == labels).sum().item()
117                total_val += batch_size

118
119            val_loader_tqdm.set_postfix({
120                'val_loss': f'{val_loss/total_val:.4f}',
121                'val_acc': f'{correct_val/total_val:.4f}'
122            })

123        val_epoch_loss = val_loss / len(val_loader.dataset)
124        val_epoch_acc = correct_val / total_val

125
126        print(f"\nEpoch {epoch+1}/{epochs} Complete")

```

```

127         print(f"Train Loss: {epoch_loss:.4f}, Train Acc: {epoch_acc
128             :.4f}")
129         print(f"Val    Loss: {val_epoch_loss:.4f}, Val    Acc: {
130             val_epoch_acc:.4f}")
131         print("-----")
132
133 if __name__ == "__main__":
134     batch_size = 32
135     train_loader = DataLoader(train_dataset, batch_size=batch_size,
136     shuffle=True)
137     val_loader = DataLoader(val_dataset, batch_size=batch_size)
138
139     device = torch.device("cuda" if torch.cuda.is_available() else "cpu"
140 )
141
142     original_densenet = models.densenet121(pretrained=True)
143     base_model = FineTunedDenseNet(pretrained_model=original_densenet,
144     num_classes=4)
145     base_model.load_state_dict(torch.load("/Users/arthe/honors/
146     best_densenet_model.pth", map_location=device))
147     base_model = base_model.to(device)
148
149     model = AnomalyDenseNet(base_model).to(device)
150
151     model.train_model("anomaly_model", batch_size, 2, 0.001,
152     train_loader, val_loader, device=device)
153
154     model.to('cpu')
155     torch.save(model.state_dict(), "anomaly_detection_model.pth")
156
157     print("\n Checking prediction scores on normal (clean) images...")
158     model.eval()
159     for i in range(5):
160         img, _ = original_binary[i]
161         input_tensor = img.unsqueeze(0)
162         with torch.no_grad():
163             raw_output = model(input_tensor)
164             score = torch.sigmoid(raw_output).item()
165             print(f"[Clean] Image {i+1} Score: {score:.4f} (Expected:
166             ~0)")
167
168     print("\nChecking prediction scores on adversarial images...")
169     for i in range(5):
170         img, _ = adversarial_binary[i]
171         input_tensor = img.unsqueeze(0)
172         with torch.no_grad():
173             raw_output = model(input_tensor)
174             score = torch.sigmoid(raw_output).item()

```

```
167     print(f"[Adversarial] Image {i+1} Score: {score:.4f} (\n    Expected: ~1)")
```

Listing 4: anomaly.py

D Appendix D: Web Application Deployment

This section of code entails the creation of a local web application to test the original image classification model, attack implementations, and anomaly detection model. A Gemini API was utilized to create a chatbot function, demonstrating the use of AI in improving patient education.

```
1 from flask import Flask, request, render_template, jsonify,
2     send_from_directory
3 from werkzeug.utils import secure_filename
4 import os
5 import time
6 import torchvision
7 from model_handler import MRI_Classifier
8 import logging
9 import traceback
10 import io
11 from PIL import Image
12 import json
13 import sys
14 import google.generativeai as genai
15 from dotenv import load_dotenv
16 load_dotenv()
17 genai.configure(api_key=os.getenv("GEMINI_API_KEY"))
18
19 app = Flask(__name__)
20 app.config['UPLOAD_FOLDER'] = 'uploads'
21 app.config['ALLOWED_EXTENSIONS'] = {'png', 'jpg', 'jpeg'}
22 app.config['MAX_CONTENT_LENGTH'] = 16 * 1024 * 1024
23
24 logging.basicConfig(level=logging.INFO)
25 logger = logging.getLogger(__name__)
26
27 model = MRI_Classifier('best_densenet_model.pth')
28
29 PROJECT_ID = "mriclassifier"
30 REGION = "us-central1"
31 gemini_model = genai.GenerativeModel("gemini-1.5-flash")
32 gemini_chat = gemini_model.start_chat(history[])
33
34
35 def allowed_file(filename):
```

```

36     return '.' in filename and \
37         filename.rsplit('.', 1)[1].lower() in app.config['
38             ALLOWED_EXTENSIONS']
39
39 def get_gemini_diagnostic(image_bytes, classification_label,
40     is_adversarial, anomaly_score):
41     try:
42         genai.configure(api_key=os.getenv("GEMINI_API_KEY"))
43         image = Image.open(io.BytesIO(image_bytes)).convert("RGB")
44
44         prompt = [
45             image,
46             f"This is an MRI brain tumor image. The classification
47             result is: {classification_label}.",
47             f"Adversarial status: {'Potentially adversarial' if
48             is_adversarial else 'Likely not adversarial'}.",
48             f"Anomaly score (if adversarial): {anomaly_score:.4f}." ,
49             "Provide a concise diagnostic explanation in JSON format
50             with ONLY this exact structure:",
50             "{",
51             '"key_findings": "...",',
51             '"potential_implications": "...",',
51             '"recommended_next_steps": "...",',
51             "}",
52             "Return ONLY the JSON object with no additional text,
53             explanations, or markdown formatting."
53         ]
54
55         response = gemini_chat.send_message(prompt)
56
57     try:
58         return json.loads(response.text)
59     except json.JSONDecodeError:
60         start_idx = response.text.find('{')
61         end_idx = response.text.rfind('}') + 1
62         if start_idx != -1 and end_idx != -1:
63             json_str = response.text[start_idx:end_idx]
64             return json.loads(json_str)
65         else:
66             return {
67                 "error": "Could not extract JSON from response",
68                 "raw_response": response.text
69             }
70     except Exception as e:
71         logger.error(f"Error in Gemini diagnostic: {str(e)}")
72         return {
73             "error": "Error generating diagnostic",
74             "exception": str(e)
75

```

```

78     }
79
80 @app.route('/')
81 def home():
82     return render_template('index.html')
83 @app.route('/chat', methods=['POST'])
84 def chat_with_gemini():
85     user_message = request.json.get('message')
86     try:
87         response = gemini_chat.send_message(user_message)
88         return jsonify({'response': response.text})
89     except Exception as e:
90         logger.error("Error in Gemini chat:")
91         logger.error(traceback.format_exc())
92         return jsonify({'error': 'Unable to get response from Gemini'})
93     , 500
94
95 @app.route('/predict', methods=['POST'])
96 def predict_upload():
97     if 'file' not in request.files:
98         return jsonify({'error': 'No file uploaded'}), 400
99
100    file = request.files['file']
101
102    if file.filename == '':
103        return jsonify({'error': 'No selected file'}), 400
104
105    if not allowed_file(file.filename):
106        return jsonify({'error': 'File type not allowed'}), 400
107
108    try:
109        os.makedirs(app.config['UPLOAD_FOLDER'], exist_ok=True)
110        timestamp = str(int(time.time()))
111        filename = secure_filename(f"{timestamp}_{file.filename}")
112        filepath = os.path.join(app.config['UPLOAD_FOLDER'], filename)
113        file.save(filepath)
114
115        result = model.predict(filepath)
116        print(f"[DEBUG] Flask response anomaly status: {result['anomaly']}", file=sys.stdout, flush=True)
117
118        display_path = os.path.join(app.config['UPLOAD_FOLDER'], f"display_{filename}")
119        torchvision.utils.save_image(result['display_image'],
120                                     display_path)
121
122        probabilities = {model.class_names[i]: f"{prob*100:.2f}%"}

```

```

121         for i, prob in enumerate(result['probabilities']
122             ])
123
124             with open(filepath, "rb") as image_file:
125                 image_bytes = image_file.read()
126
127             diagnostic_info = get_gemini_diagnostic(
128                 image_bytes,
129                 result['class'],
130                 result['anomaly']['is_anomalous'],
131                 result['anomaly']['anomaly_score']
132             )
133
134             return jsonify({
135                 'status': 'success',
136                 'prediction': result['class'],
137                 'confidence': f"{result['confidence']*100:.2f}%",
138                 'probabilities': probabilities,
139                 'display_image': f"/display/{filename}",
140                 'is_adversarial': result['anomaly']['is_anomalous'],
141                 'anomaly_score': float(result['anomaly']['anomaly_score']),
142                 'diagnostic_info': diagnostic_info
143             })
144
145     except Exception as e:
146         logger.error(f"Error: {str(e)}")
147         logger.error(traceback.format_exc())
148         return jsonify({'error': str(e)}), 500
149     finally:
150         if 'filepath' in locals() and os.path.exists(filepath):
151             os.remove(filepath)
152
153 @app.route('/display/<filename>')
154 def serve_display_image(filename):
155     return send_from_directory(app.config['UPLOAD_FOLDER'], f"display_{filename}")
156
157 if __name__ == '__main__':
158     os.makedirs(app.config['UPLOAD_FOLDER'], exist_ok=True)
159     app.run(host='0.0.0.0', port=5004, debug=True)

```

Listing 5: app.py

```

1 <!DOCTYPE html>
2 <html lang="en">
3 <head>
4     <meta charset="UTF-8">
5     <meta name="viewport" content="width=device-width, initial-scale
=1.0">

```

```

6   <title>MRI Imaging Classifier</title>
7   <link href="https://cdn.jsdelivr.net/npm/bootstrap@5.1.3/dist/css/
8     bootstrap.min.css" rel="stylesheet">
9   <link rel="stylesheet" href="https://cdn.jsdelivr.net/npm/bootstrap-
10    -icons@1.11.3/font/bootstrap-icons.min.css">
11   <script src="https://cdn.jsdelivr.net/npm/bootstrap@5.1.3/dist/js/
12     bootstrap.bundle.min.js"></script>
13   <script src="https://cdn.jsdelivr.net/npm/marked/marked.min.js"></
14     script>
15
16   <style>
17     .adversarial-alert {
18       animation: pulse 2s infinite;
19     }
20     @keyframes pulse {
21       0% { box-shadow: 0 0 0 0 rgba(220, 53, 69, 0.7); }
22       70% { box-shadow: 0 0 0 10px rgba(220, 53, 69, 0); }
23       100% { box-shadow: 0 0 0 0 rgba(220, 53, 69, 0); }
24     }
25     .security-meter {
26       height: 20px;
27       background: linear-gradient(to right, #28a745, #ffc107, #
28         dc3545);
29       border-radius: 10px;
30       margin-bottom: 15px;
31     }
32     .security-indicator {
33       height: 100%;
34       width: 0%;
35       background-color: rgba(255, 255, 255, 0.7);
36       border-radius: 10px;
37       transition: width 1s;
38     }
39     body {
40       background-color: #f8f9fa;
41       font-family: 'Segoe UI', Tahoma, Geneva, Verdana, sans-
42       serif;
43     }
44     .container {
45       max-width: 800px;
46     }
47     .upload-area {
48       border: 2px dashed #6c757d;
49       border-radius: 10px;
50       padding: 2rem;
51       text-align: center;
52       cursor: pointer;
53       transition: all 0.3s;

```

```

48         background-color: white;
49         margin-bottom: 2rem;
50     }
51     .upload-area:hover {
52         border-color: #0d6efd;
53         background-color: #f0f7ff;
54     }
55     #previewImage {
56         max-width: 100%;
57         max-height: 300px;
58         border-radius: 8px;
59         display: none;
60         margin: 0 auto 1.5rem auto;
61     }
62     .result-card {
63         background-color: white;
64         border-radius: 10px;
65         padding: 1.5rem;
66         box-shadow: 0 4px 6px rgba(0, 0, 0, 0.1);
67         margin-top: 1.5rem;
68         display: none;
69     }
70
71     .diagnosis-badge {
72         font-size: 1.2rem;
73         padding: 0.5rem 1rem;
74         border-radius: 5px;
75     }
76     .diagnostic-card {
77         background-color: #e9ecf;
78         border-radius: 10px;
79         padding: 1.5rem;
80         margin-top: 1.5rem;
81         box-shadow: 0 2px 4px rgba(0, 0, 0, 0.05);
82         display: none;
83     }
84     .diagnostic-card h5 {
85         color: #495057;
86         margin-bottom: 1rem;
87     }
88     .diagnostic-card p {
89         color: #6c757d;
90         line-height: 1.6;
91     }
92     .glioma { background-color: #fff3cd; color: #856404; }
93     .meningioma { background-color: #d4edda; color: #155724; }
94     .notumor { background-color: #cce5ff; color: #004085; }
95     .pituitary { background-color: #f8d7da; color: #721c24; }

```

```

96     </style>
97 </head>
98 <body>
99     <div class="container py-5">
100         <div class="text-center mb-5">
101             <h1 class="display-4">Brain Tumor MRI Classifier</h1>
102             <p class="lead">Upload an MRI scan to classify its tumor
103                 type!</p>
104             </div>
105
106             <div id="uploadArea" class="upload-area">
107                 <h4><i class="bi bi-cloud-arrow-up"></i> Upload MRI Image</
108                 h4>
109                 <p class="text-muted">Click to browse or drag and drop</p>
110                 <input type="file" id="fileInput" accept="image/*" class="d
111                 -none">
112             </div>
113
114
115             <div class="text-center">
116                 <img id="previewImage" class="img-fluid">
117             </div>
118
119             <button id="predictBtn" class="btn btn-primary btn-lg w-100 py
120                 -3" disabled onclick="predict()">
121                 Classify Image
122             </button>
123             <div id="result" class="result-card" style="display: none;">
124                 <h3>Diagnosis: <span id="predictionResult"></span></h3>
125                 <p>Confidence: <span id="confidenceValue"></span></p>
126
127                 <div id="adversarialAlert" class="alert alert-danger" style
128                     ="display: none;">
129                     <i class="bi bi-shield-exclamation"></i> Warning:
130                     Potential adversarial image detected!
131                     <div class="mt-2">
132                         <div class="security-meter">
133                             <div id="securityIndicator" class="security-
134                             indicator"></div>
135                         </div>
136                     </div>
137             </div>
138
139             <h5 class="mt-4">Probability Distribution:</h5>
140             <div id="probabilityBars"></div>
141         </div>
142         <div id="diagnosticInfo" class="diagnostic-card">
143             <h5>Diagnostic Information:</h5>
144             <p id="diagnosticText"></p>

```

```

137     </div>
138
139     <div id="chat-container" style="margin-top: 30px;">
140         <h3>Follow-up Chat with Gemini</h3>
141         <div id="chat-log" style="border: 1px solid #ccc; padding: 10px; height: 200px; overflow-y: auto; background: #f9f9f9; border-radius: 8px;"></div>
142         <div style="margin-top: 10px;">
143             <input type="text" id="chat-input" placeholder="Ask a follow-up question..." style="width: 80%; padding: 8px;" />
144             <button onclick="sendChat()" style="padding: 8px 12px;">Send</button>
145         </div>
146     </div>
147 </div>
148
149 <script>
150     const uploadArea = document.getElementById('uploadArea');
151     const fileInput = document.getElementById('fileInput');
152     const previewImage = document.getElementById('previewImage');
153     const predictBtn = document.getElementById('predictBtn');
154     const resultDiv = document.getElementById('result');
155
156     uploadArea.addEventListener('dragover', (e) => {
157         e.preventDefault();
158         uploadArea.style.borderColor = '#0d6efd';
159         uploadArea.style.backgroundColor = '#f0f7ff';
160     });
161
162     uploadArea.addEventListener('dragleave', () => {
163         uploadArea.style.borderColor = '#6c757d';
164         uploadArea.style.backgroundColor = 'white';
165     });
166
167     uploadArea.addEventListener('drop', (e) => {
168         e.preventDefault();
169         uploadArea.style.borderColor = '#6c757d';
170         uploadArea.style.backgroundColor = 'white';
171         if (e.dataTransfer.files.length) {
172             fileInput.files = e.dataTransfer.files;
173             previewFile(fileInput.files[0]);
174         }
175     });
176
177     uploadArea.addEventListener('click', () => {
178         fileInput.click();
179     });
180

```

```

181     function previewFile(file) {
182         const reader = new FileReader();
183         reader.onload = (e) => {
184             previewImage.src = e.target.result;
185             previewImage.style.display = 'block';
186             predictBtn.disabled = false;
187             resultDiv.style.display = 'none';
188         };
189         reader.readAsDataURL(file);
190     }
191
192     function predict() {
193         if (!fileInput.files.length) {
194             showAlert('Please select an image first', 'warning');
195             return;
196         }
197
198         const formData = new FormData();
199         formData.append('file', fileInput.files[0]);
200
201         predictBtn.disabled = true;
202         predictBtn.innerHTML =
203             '<span class="spinner-border spinner-border-sm" role="status" aria-hidden="true"></span>
204             Processing...
205             ';
206
207         resultDiv.style.display = 'none';
208
209         fetch('/predict', {
210             method: 'POST',
211             body: formData
212         })
213         .then(response => {
214             if (!response.ok) {
215                 throw new Error(`HTTP error! status: ${response.
216 status}`);
217             }
218             return response.json();
219         })
220         .then(data => {
221             console.log("Probabilities:", data.probabilities);
222
223             displayResults(data);
224             data.probabilities = {
225                 benign: 0.75,
226                 early: 0.1,
227                 pre: 0.1,

```

```

227         pro: 0.05
228     };
229     predictBtn.disabled = false;
230     predictBtn.innerHTML = 'Classify Image';
231   })
232   .catch(error => {
233     console.error('Error:', error);
234     showAlert('An error occurred while processing the image
235   ', 'danger');
236
237   predictBtn.disabled = false;
238   predictBtn.innerHTML = 'Classify Image';
239 });
240
241 function showAlert(message, type) {
242   const existingAlert = document.getElementById('dynamicAlert
243 ');
244   if (existingAlert) {
245     existingAlert.remove();
246   }
247
248   const alertDiv = document.createElement('div');
249   alertDiv.id = 'dynamicAlert';
250   alertDiv.className = `alert alert-${type} alert-dismissible
251 fade show`;
252   alertDiv.setAttribute('role', 'alert');
253   alertDiv.innerHTML =
254     `${message}
255       <button type="button" class="btn-close" data-bs-dismiss
256 ="alert" aria-label="Close"></button>
257     `;
258
259   uploadArea.after(alertDiv);
260
261   setTimeout(() => {
262     const alert = bootstrap.Alert.getInstance(alertDiv);
263     if (alert) {
264       alert.close();
265     }
266   }, 5000);
267 }
268
269 function displayResults(data) {
270   console.log("Received data from backend:", data);
271
272   const predictionResult = document.getElementById('
273 predictionResult');
```

```

270     const advAlert = document.getElementById('adversarialAlert');
271     const diagnosticDiv = document.getElementById('diagnosticInfo');
272     const diagnosticText = document.getElementById('diagnosticText');
273     document.getElementById("predictionResult").textContent = data.prediction;
274
275
276     predictionResult.textContent = data.prediction;
277     predictionResult.className = 'diagnosis-badge' + data.
278     prediction.toLowerCase();
279     document.getElementById('confidenceValue').textContent =
280     data.confidence;
281
282
283     const probabilityBars = document.getElementById('probabilityBars');
284     probabilityBars.innerHTML = '';
285
286
287     for (const [className, prob] of Object.entries(data.
288     probabilities)) {
289         const barContainer = document.createElement('div');
290         barContainer.className = 'mb-3';
291
292         const labelRow = document.createElement('div');
293         labelRow.className = 'd-flex justify-content-between mb-
294         -1';
295
296         const label = document.createElement('span');
297         label.textContent = className;
298         label.className = 'text-capitalize fw-medium';
299
300         const percentage = document.createElement('span');
301         percentage.textContent = prob;
302
303         labelRow.appendChild(label);
304         labelRow.appendChild(percentage);
305         barContainer.appendChild(labelRow);
306
307         const progress = document.createElement('div');
308         progress.className = 'probability-bar';
309
310         const fill = document.createElement('div');
311         fill.className = 'probability-fill';
312         fill.style.width = `${Math.min(100, (prob * 100).
313        toFixed(1))}%`; // Cap at 100%

```

```

308         progress.appendChild(fill);
309         barContainer.appendChild(progress);
310         probabilityBars.appendChild(barContainer);
311     }
312     document.getElementById("result").style.display = "block";
313     if (data.is_adversarial) {
314         advAlert.style.display = 'block';
315         advAlert.classList.add('adversarial-alert');
316     } else {
317         advAlert.style.display = 'none';
318     }
319
320     if (data.diagnostic_info) {
321         if (data.diagnostic_info.error) {
322             diagnosticText.innerHTML =
323                 '<div class="alert alert-warning">
324                 <strong>Note:</strong> ${data.
325                 diagnostic_info.error}
326                 ${data.diagnostic_info.raw_response} ?
327                 <div class="mt-2"><small>${data.
328                 diagnostic_info.raw_response}</small></div>' : ''
329             '
```

- 329 } else {
- 330 diagnosticText.innerHTML = formatDiagnosticInfo(
- 331 data.diagnostic_info, data.anomaly_score);
- 332 }
- 333 diagnosticDiv.style.display = 'block';
- 334 } else {
- 335 diagnosticText.innerHTML = '<p>No diagnostic information available</p>';
- 336 diagnosticDiv.style.display = 'block';
- 337 }
- 338 }
- 339
- 340 function formatDiagnosticInfo(info, anomaly_score) {
- 341 if (!info) return '<p>No diagnostic information available</p>';
- 342
- 343 let html = '';
- 344
- 345 if (parseFloat(anomaly_score) > 0.9) {
- 346 html += '
- 347 <div class="alert alert-danger mb-3">
- 348 Important Note: This scan
- 349 shows a high anomaly score (\${anomaly_score}),

```

350             indicating significant uncertainty in the
351             classification. Expert review is strongly recommended.
352         </div>
353     `;
354 }
355
356 if (info.key_findings) {
357     html += `
358         <h5 class="mt-3">Key Findings</h5>
359         <p>${info.key_findings}</p>
360     `;
361 }
362
363 if (info.potential_implications) {
364     html += `
365         <h5 class="mt-3">Potential Implications</h5>
366         <p>${info.potential_implications}</p>
367     `;
368 }
369
370 if (info.recommended_next_steps) {
371     html += `
372         <h5 class="mt-3">Recommended Next Steps</h5>
373         <p>${info.recommended_next_steps}</p>
374     `;
375 }
376
377 html += `
378     <div class="alert alert-secondary mt-3">
379         <strong>Disclaimer:</strong> This AI analysis is
for informational purposes only and
does not constitute medical advice. Always consult
with a qualified healthcare
professional for diagnosis and treatment.
380     </div>
381 `;
382
383 return html;
384 }
385
386 function sendChat() {
387     const input = document.getElementById("chat-input");
388     const message = input.value;
389     if (!message.trim()) return;
390
391     appendChatMessage("You", message);
392     input.value = "";
393
394     fetch("/chat", {

```

```

395         method: "POST",
396         headers: {
397             "Content-Type": "application/json"
398         },
399         body: JSON.stringify({ message: message })
400     })
401     .then(response => response.json())
402     .then(data => {
403         if (data.response) {
404             appendChatMessage("Gemini", data.response);
405         } else {
406             appendChatMessage("Gemini", "Sorry, I couldn't
process that.");
407         }
408     })
409     .catch(err => {
410         appendChatMessage("Gemini", "Error connecting to server
.");
411         console.error(err);
412     });
413 }

414
415     async function appendChatMessage(sender, message) {
416         const chatLog = document.getElementById("chat-log");
417         const newMessage = document.createElement("div");
418         newMessage.innerHTML = `<strong>${sender}</strong> <p></p>`;
419         chatLog.appendChild(newMessage);
420
421         if (sender === "Gemini") {
422             const parsedMessage = marked.parse(message);
423             const p = newMessage.querySelector("p");
424
425             let i = 0;
426             const tempDiv = document.createElement("div");
427             tempDiv.innerHTML = parsedMessage;
428             const text = tempDiv.textContent || tempDiv.innerText ||
"";
429
430             const typingSpeed = 10;
431
432             function typeChar() {
433                 if (i < text.length) {
434                     p.textContent += text[i++];
435                     chatLog.scrollTo({ top: chatLog.scrollHeight,
behavior: "smooth" });
436                     setTimeout(typeChar, typingSpeed);
437                 } else {
438                     p.innerHTML = parsedMessage;


```

```
439         }
440     }
441
442     typeChar();
443 } else {
444     newMessage.innerHTML = `<strong>${sender}</strong> <p>${
445 message}</p>`;
446 }
447 </script>
448 </body>
449 </html>
```

Listing 6: index.html

Adversarial Machine Learning

Arthea Valderrama



Agenda

01

Overview

02

Machine
Learning

03

Carlini &
Wagner

04

SimBA

05

Anomaly
Detection


06

Results




Motivation

- Machine Learning (ML) is widely used
- Potential to improve quality of life
- However, it is vulnerable to attacks



What is Machine Learning?



**Machine
learning**

Input data

Images of MRI scans

Relationships

Patterns


Dependencies

Structures

Output

Categorized tumors

Further applications in Healthcare



Disease surveillance

Identify spikes in hospital visits

Assist clinicians

Reduce clinician workload

Enhance treatment


Personalized treatment plans

Patient Education

Accessible to health literacy

Adversarial Attacks


- Malicious inputs designed to deceive models
- Imperceptible to the human eye
- Cause mass misdiagnoses, hurt treatment plans
- Incentives- exploit hospitals' finances, cause chaos




Carlini & Wagner

- Highly effective
- Requires model information
- Optimized math
- Consistent in success


Original: pituitary




Adversarial: meningioma



Original: notumor




Adversarial: meningioma




SimBA

- Does not need model information
- Can be more perceptible to the human eye
- Requires more computational resources
- Easier to implement

notumor




meningioma




Anomaly Detection


- Identify deviation from normal data patterns
- Cover weaknesses of image classification
- Flag adversarial examples



Deployment Pipeline




Demo




Results

- higher success rate, but more visually **distorted**
- C&W generally more successful
- SimBA much more faster
- Anomaly detection strong in attacks





Implications

- Account for different types of attacks
- Potential to decrease reliability of AI
- Generalizable- can affect different AI



Conclusion

- Important to incorporate adversarial samples in ML
- ML in healthcare is promising but vulnerable
- Continue to strengthen defense mechanisms while utilizing AI in day-to-day life



Thanks!

Do you have any questions?



CREDITS: This presentation template was created by [Slidesgo](#), and includes icons by [Flaticon](#), and infographics & images by [Freepik](#)

Please keep this slide for attribution

References

- [1] N. Carlini and D. Wagner, [Towards evaluating the robustness of neural networks](#) (2017), arXiv:1608.04644 [cs.CR].
- [2] C. Guo, J. R. Gardner, Y. You, A. G. Wilson, and K. Q. Weinberger, [Simple black-box adversarial attacks](#) (2019), arXiv:1905.07121 [cs.LG].
- [3] T. R. Sarkar, N. Das, P. S. Maitra, B. Some, R. Saha, O. Adhikary, B. Bose, and J. Sen, [Evaluating adversarial robustness: A comparison of fgsm, carlini-wagner attacks, and the role of distillation as defense mechanism](#) (2024), arXiv:2404.04245 [cs.CR].
- [4] M. Nickparvar, [Brain tumor mri dataset](#) (2021).
- [5] Carco-git, Cw_attack_on_mnist, https://github.com/Carco-git/CW_Attack_on_MNIST (2023), accessed: 2025-04-21.
- [6] cg563, Simple_blackbox-attack, <https://github.com/cg563/simple-blackbox-attack/tree/master> (2023), accessed: 2025-04-21.